

1 Data Mining: BeeViva Challenges

2
3 MARCO BELLÒ, University of Padua, Italy
4

5 All code, images, L^AT_EX sources, as well as a less polished version of this report in markdown format, split into two files for the two
6 challenges, can be found at the following repository: <https://github.com/mhetacc/DataMiningChallenges/>.

7 To be precise:

- 8 • Rice Varieties markdown report: https://github.com/mhetacc/DataMiningChallenges/blob/main/RiceChallenge/Rice_report.md
- 9 • Rice Varieties R code: https://github.com/mhetacc/DataMiningChallenges/blob/main/RiceChallenge/rice_challenge.r
- 10 • Phone Users markdown report: https://github.com/mhetacc/DataMiningChallenges/blob/main/PhoneUserChallenge/Phone_report.md
- 11 • Phone Users R code: https://github.com/mhetacc/DataMiningChallenges/blob/main/PhoneUserChallenge/phone_challenge.r

16 ACM Reference Format:

17 Marco Bellò. 2025. Data Mining: BeeViva Challenges. 1, 1 (December 2025), ?? pages. <https://doi.org/XXXXXXX.XXXXXXX>

19 1 RICE VARIETIES

21 1.1 Preliminary Observations

23 1.1.1 *Dataset.* All following considerations are made using the datasets provided on the challenge page (*rice_test.csv*
24 and *rice_train.csv*), but it is worth noting that they both stem from an original one that can be found at the [dataset](#)
25 [source page](#). I used it to compute an alternative version of the *test* dataset that contains the true attribute for *Class*.

27 Listing 1. Merge datasets

```
28 df1 = pd.read_csv("rice_test.csv").round(10)
29 df2 = pd.read_csv("Rice_Cammeo_Osmancik.csv").round(10)
30
31 keys = [col for col in df1.columns]
32 merged = df1.merge(df2, on=keys, how="inner")
```

34 1.1.2 *Scatter Plot.* From the *scatter plot* we can infer that there are four features (*Area*, *Perimeter*, *Convex_Area* and
35 *Major_Axis_Length*) that seem to be extremely correlated

37 1.1.3 *Correlation Matrix.* We can use a correlation matrix to see exactly how related to each other the features are.
38 The results are as follow.

40 As suspected, all four features mentioned above have a high degree of correlation (over ninety percent). Once way to
41 handle this is to either use models robust against collinearity, or to either combine or remove some of the correlated
42 features.

44 Author's address: Marco Bellò, marco.bello.3@studenti.unipd.it, University of Padua, Padua, Italy.

45 Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not
46 made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components
47 of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on
48 servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

49 © 2025 Copyright held by the owner/author(s). Publication rights licensed to ACM.

50 Manuscript submitted to ACM

52 Manuscript submitted to ACM

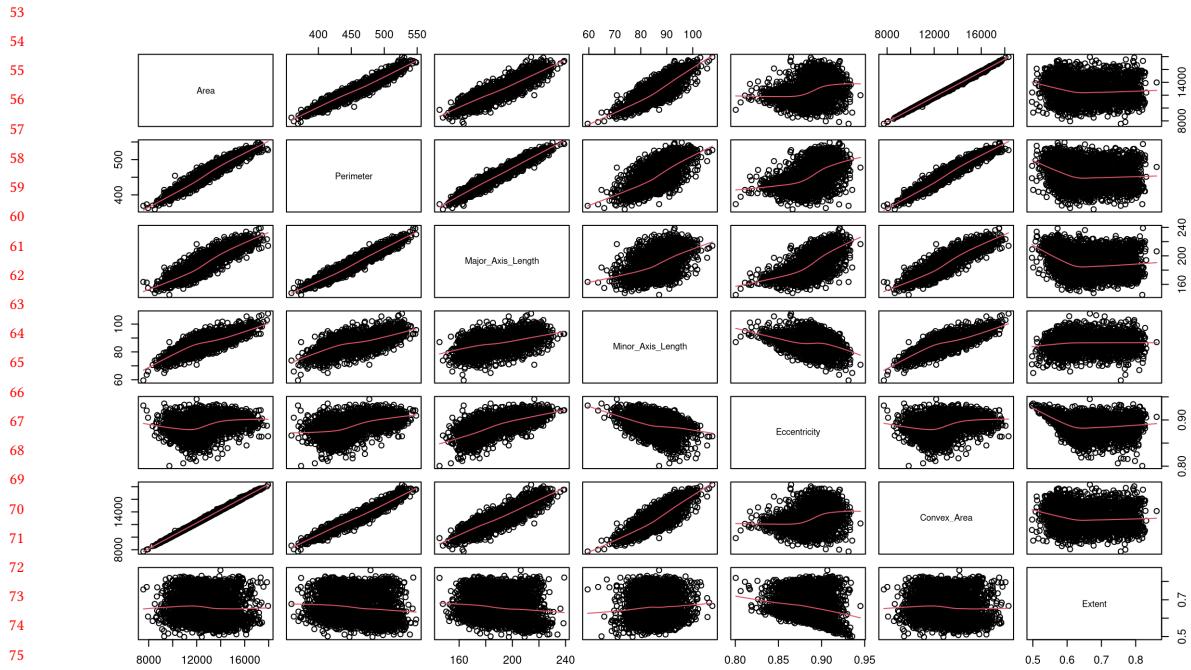


Fig. 1. Scatter Plot for the *rice_train* dataset with LOESS smoothing lines for each class.

Table 1. Correlation matrix of morphological parameters

Area	Perimeter	Major_Axis_Length	Minor_Axis_Length	Eccentricity	Convex_Area	Extent	
Area	1.00000000	0.96704011	0.90356325	0.79022701	0.34653177	0.99895272	- 0.06002223
Perimeter	0.96704011	1.00000000	0.97163180	0.63486482	0.53784650	0.97046788	- 0.12718015
Major_Axis_Length	0.90356325	0.97163180	1.00000000	0.45675159	0.70625660	0.90390912	- 0.13562592
Minor_Axis_Length	0.79022701	0.63486482	0.45675159	1.00000000	- 0.29393931	0.78975480	0.06192558
Eccentricity	0.34653177	0.53784650	0.70625660	- 0.29393931	1.00000000	0.34707340	- 0.19301157
Convex_Area	0.99895272	0.97046788	0.90390912	0.78975480	0.34707340	1.00000000	- 0.06397213
Extent	- 0.06002223	- 0.12718015	- 0.13562592	0.06192558	- 0.19301157	- 0.06397213	1.00000000

1.1.4 *Features' Importance*. To get a better idea on the importance of each feature, I trained a simple linear regression model, and looked at the result with `summary(fit)`. The results follows.

Table 2. Regression coefficients for morphological parameters

Parameter	Estimate	Std. Error	t value	Pr(> t)	Importance
(Intercept)	-2.152e+00	1.990e+00	-1.081	0.279633	
Area	5.601e-04	1.023e-04	5.474	4.79e-08	***
Perimeter	8.440e-03	2.221e-03	3.800	0.000148	***
Major_Axis_Length	-2.198e-02	5.709e-03	-3.851	0.000120	***
Minor_Axis_Length	4.669e-02	1.213e-02	3.850	0.000121	***
Eccentricity	4.534e+00	1.828e+00	2.481	0.013170	*
Convex_Area	-8.609e-04	9.418e-05	-9.141	< 2e-16	***
Extent	7.146e-02	7.144e-02	1.000	0.317237	

The significance stars tells us that, with the exception of *Minor_Axis_Length*, the features that are highly correlated to each other contributes strongly to the model, while the *Eccentricity* and *Extent* contribute very little.

We can also measure the total variance explained by all predictors combined using $R^2 = 0.6953849$.

2 PYTHON

We are now going to discuss the technologies employed in the development of the project, as well as presenting our implementation of the Raft's algorithm.

Python is a high-level, dynamically typed and interpreted programming language that is often used for scripting, data analysis and small application development, making it a non-obvious choice for this project, which does not fall into any of these categories.

As a language, it has two main advantages compared to others: first of all it is undoubtedly the most popular and widely used in the world (figure ??) [? ?], meaning abundant documentation and resources. Secondly it has a huge ecosystem of libraries that implement all the functionalities we need for this project, namely: *XML-RPC* for the remote procedure calls (RPCs), *threading* to handle local concurrency and *Pygame* to manage everything game-related.

2.1 Remote Procedure Calls

In Raft's specifications it is stated that nodes communicate with each other via remote procedure calls [?], which in distributed computing is when a program causes a procedure (or subroutine) to execute in another address space (commonly on another computer on a shared network) calling it as if it were local (that is, the programmer writes the same code whether the subroutine is local or remote).

There are many libraries that implement this functionality, like gRPC (<https://grpc.io/>), which is a high performance open source RPC framework used by many big players, such as Netflix ¹ and Cockroach Labs ², available for many languages (Python included), but we opted for the standard library *XML-RPC* ³ thanks to its promised simplicity and ease of use.

The library provides both server and client implementations, encapsulating the former in its own loop, while the latter can be fired as needed allowing a bit more flexibility in its usage.

¹Netflix Ribbon is an Inter Process Communication library built in software load balancers: <https://github.com/Netflix/ribbon>

²Cockroach Labs is the company behind CockroachDB, a highly resilient distributed database: <https://www.cockroachlabs.com/>

³XML-RPC is a Remote Procedure Call method that uses XML passed via HTTP as a transport: <https://docs.python.org/3/library/xmlrpc.html>

```

157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176 images/TIOBEindex.png
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208

```

Fig. 2. TIOBE Programming Community Index, focus on Python statistics, 2025. (<https://www.tiobe.com/tiobe-index/>)

In code ??, *client* is an instance of *ServerProxy*, which acts as the client-side interface for XML-RPC, allowing it to call the remote procedure *test_foo* as if it were a local function, even though it executes on a server in a different networked location.

Listing 2. Client as server proxy

```

1 with xmlrpclib.ServerProxy('http://localhost:8000', allow_none=True) as client:
2     print(client.test_foo(42)) # print returned value

```

209 The server must be instantiated and kept running by calling its event loop (e.g., using `serve_forever`), and all remote
 210 procedure calls must be registered using the `register_function` method of `SimpleXMLRPCServer` (code ??).

212 Listing 3. Server

```
213
214 1 with SimpleXMLRPCServer (('localhost', 8000)) as server:
215 2     def test_foo(number):
216 3         return f'The number is {number}'
217 4
218 5     server.register_function(test_foo)
219 6 server.serve_forever() # keep server alive
```

220 For this project, we extended `SimpleXMLRPCServer` to create a class that implements the Raft protocol (more details
 221 in section ??).

223 2.2 Concurrency

224 In this project the need for concurrent programming arises from two challenges: every server has an internal timer that
 225 fires at certain intervals, and every node has to run a game engine and the server itself at the same time, both of which
 226 are, by design of their own respective libraries, independent, blocking and perpetually executing loops.

227 Some Raft implementations we examined achieve concurrency through asynchronous programming⁴, using libraries
 228 such as `asyncio`⁵, thereby avoiding the need to manage common multithreading challenges like ensuring thread-safety
 229 by preventing race conditions or data corruption and other hazards such as lock-free reordering or incorrect granularity
 230 (see footnote for more⁶). That being said, while powerful and efficient, writing asynchronous code can be awkward and
 231 cumbersome, so we opted for a more traditional approach using multithreaded programming: in computer science, a
 232 thread of execution is the smallest sequence of programmed instructions that can be scheduled independently [?], and
 233 multiple threads may be executed concurrently sharing resources such as memory. This is directly counterpointed to
 234 multiprocessing, where each process has its own storage space, and moreover processes are typically made of threads.
 235 Processes and threads are profoundly different and do not serve the same purpose, but it is useful to cite both of them
 236 to provide the context needed to fully understand section ??.

237 In Python there are modules in the standard library for both of them, respectively `threading`⁷ and `multiprocessing`⁸.
 238 It is fundamental to note that the former does not provide real multi-threading since, due to the Global Interpreter Lock
 239 of CPython (the, for want of a better word, official Python implementation), only one thread can execute bytecode at
 240 once. To cite directly from the documentation: "[GIL is] The mechanism used by the CPython interpreter to assure that
 241 only one thread executes Python bytecode at a time. This simplifies the CPython implementation by making the object model
 242 (including critical built-in types such as dict) implicitly safe against concurrent access. Locking the entire interpreter makes
 243 it easier for the interpreter to be multi-threaded, at the expense of much of the parallelism afforded by multi-processor
 244 machines."⁹

245 ⁴Two examples of Raft's implementations that leverage asynchronous programming are Raftos (<https://github.com/zhebrak/raftos/tree/master>) and Zatt (<https://github.com/simonacca/zatt/tree/master>)

246 ⁵Asyncio is a library to write concurrent code using the `async/await` syntax: <https://docs.python.org/3/library/asyncio.html>

247 ⁶Multithreading Hazards, Microsoft: <https://learn.microsoft.com/en-us/archive/msdn-magazine/2008/october/concurrency-hazards-solving-problems-in-your-multithreaded-code>

248 ⁷The threading module provides a way to run multiple threads (smaller units of a process) concurrently within a single process: <https://docs.python.org/3/library/threading.html>

249 ⁸The multiprocessing module is a package that supports spawning processes using an API similar to the threading module: <https://docs.python.org/3/library/multiprocessing.html>

250 ⁹Global Interpreter Lock: <https://docs.python.org/3/glossary.html#term-global-interpreter-lock>

261 Thankfully, this does not apply with the *multiprocessing* module, which creates separate processes instead, offering
262 both local and remote concurrency effectively side-stepping the Global Interpreter Lock, allowing programmers to fully
263 leverage multiple cores. As previously stated, processes are much heavier than threads and thus more expensive to
264 create, but do not incur the risks of shared memory.
265

2.2.1 Comparison. To evaluate which of the two modules is more suited for our purposes, we devised a simple experiment: we created two game instances with one hundred and one thousands coloured dots respectively (figure ??), that move around by offsetting their position each frame of a random amount between minus five and plus five pixels (pseudocode ??).

Then we ran both of them in three scenarios: with the game instance alone (baselines), with a server alive in a thread and with a server alive in a process, and we measured the *frames per second* (FPS)¹⁰ in each case, since it is the most common metric to evaluate game performance. Higher FPS-count translates to a smoother and more responsive, i.e., better, gaming experience.

Listing 4. Pygame graphical dot offset

```
1 # create random offsets for both x and y coordinates
297 2 xmov = random.randint(-5,5)
298 3 ymov = random.randint(-5,5)
299
300 4
301 5 # move the dot by a certain offset
302 6 dot.move_by(xmov, ymov)
303
```

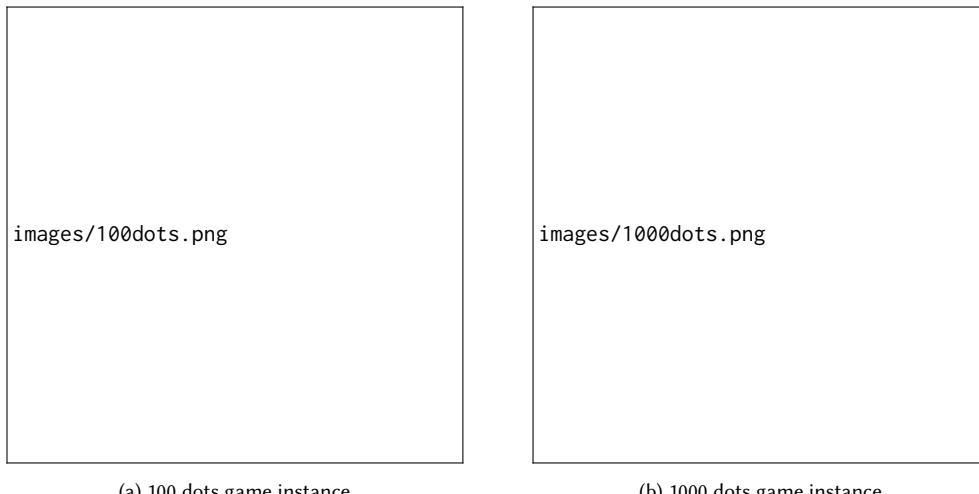


Fig. 3. Two game instances made with Pygame, with respectively 100 and 1000 dots that randomly move around.

Results, shown in the graph at figure ??, tell us us that:

- Increasing the number of dots from 100 to 1000 halves the FPS count;

³⁰⁹ ³¹⁰ ³¹¹¹⁰Frame rate, most commonly expressed in frames per second or FPS, is typically the frequency (rate) at which consecutive images (frames) are captured or displayed. This definition applies to film and video cameras, computer animation, and motion capture systems, while in the context of computer graphics is the rate at which a system, particularly the graphic card, is able to generate frames. Source: https://en.wikipedia.org/wiki/Frame_rate

- 313 • Adding a server in a separate thread halves performance;
314 • Using *multiprocessing* yields worse performance than *threading* in the 100-dots scenario (about -30%), while
315 performing similarly in the 1000-dots one.

316
317 This leads us to conclude that, for our specific purposes, the *threading* module is the best choice, especially since
318 the final game will be way less computationally expensive from a graphical standpoint, hence using a lighter weight
319 alternative should be even more beneficial than tested.
320

321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341 images/performance_eval_fps_graph.png
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360

361
362 Fig. 4. Performance evaluation graph: red hues for baselines, blue hues for threading and green hues for multiprocessing. Darker
363 shades for 1000 dots and lighter shades for 100 dots game instances

365 All tests have been performed with the following machine:

- 366
- 367 • OS: Ubuntu 24.04.1 LTS x86_64;
 - 368 • Kernel: 6.8.0-52-generic;
 - 369 • Shell: bash 5.2.21;
 - 370 • CPU: 13th Gen Intel i7-13620H;
 - 372 • GPU: NVIDIA GeForce RTX 4050 Laptop GPU;
 - 373 • Memory: 15610MiB;
 - 374 • Python version: 3.12.3;
 - 375 • Power Mode: Balanced;
 - 376 • Power Supply: 100W via type C.
- 378

379 2.3 Game Engines

380

381 There are many ways to implement a graphical user interface: from clever shell tricks like htop ¹¹, to full-fledged game
382 engines like Unity ¹² or Godot ¹³ that often come with their own editor and a *top-down* approach, meaning build the UI
383 first and then go down to code as needed for scripting and refining.

385 Unfortunately, our needs are quite opposite: what we want is a code-only, mono-language framework that while
386 slowing down game development should simplify merging Raft with it. The choice thus boiled down to two alternatives:
387 tkinter and Pygame.

389 390 2.3.1 Comparison. Let's list strengths and weaknesses of the two.

- 391
- 392 • **tkinter:**
 - 393 – Module of the standard library;
 - 394 – Few lines of code to make simple UIs;
 - 395 – Low flexibility;
 - 396 – No game loop;
 - 397 – Not a game engine;
 - 398 • **Pygame:**
 - 399 – Extreme flexibility;
 - 400 – Direct access to game loop;
 - 401 – APIs to access many kinds of user inputs;
 - 402 – Verbose to obtain simple UIs;
 - 403 – Non-standard community-made framework.
- 406

407 We ultimately decided to opt for Pygame for four reasons: it is extremely flexible, exposes many useful functions (for
408 example to catch different user inputs), gives direct access to the game loop and it is a novel and fun framework that
409 has never, to the best of our knowledge, been used in such a fashion.

413 ¹¹Htop is a cross-platform text-mode interactive process viewer: <https://htop.dev/>

414 ¹²Unity is a cross-platform game engine developed by Unity Technologies: <https://unity.com/>

415 ¹³Godot is a cross-platform open-source game engine: <https://godotengine.org/>

3 RAFT

It is not our intention to plagiarize Ongaro and Ousterhout's excellent work "*In Search of an Understandable Consensus Algorithm*" [?] by presenting the Raft algorithm's specifications. Instead, we are going to discuss how we molded it to our own use case.

The algorithm divides its nodes into three roles, namely *leader*, *follower* and *candidate*, and revolves around three core functionalities: leader election, log replication and cluster membership change. Log compaction is also mentioned, while a byzantine fault tolerant variant is never explored by the original authors. To grant consistency, Raft's design choice is to centralize all decisions on one node, the above mentioned leader, that synchronizes all cluster's nodes.

One last component, instrumental to the functioning of the algorithm, is the *term*: everything happens in a certain term, which divides time logically and increments every election. This is necessary to recognize out-of-date leaders: if some follower has a term greater than the leader's, said leader is outdated.

Our Raft class directly extends `simpleXMLRPCServer` from XML-RPC module, as shown at code ??.

Lastly, to fire off non-blocking concurrent RPCs on the cluster, we leverage the *concurrent.futures* module using *ThreadPoolExecutor*. To avoid creating and destroying pools every time a server needs to communicate with the cluster, we embedded a finite amount of workers as class attributes (code ??).

Listing 5. Class Raft definition

```
8
9     class Raft(SimpleXMLRPCServer):
10         def __init__(self,
11             addr: tuple[str, int],
12             allow_none: bool = True,
13             # ...
14             last_index_on_server: list[tuple[int, int]] | None = None
15             ):
16             SimpleXMLRPCServer.__init__(self, addr=addr, allow_none=allow_none)
```

Listing 6. ThreadPoolExecutor created with as many workers as there are servers in the cluster

```
 1 class Raft(SimpleXMLRPCServer):
 2     def __init__(self,
 3                  # ...
 4                  )
 5     # start executors pool
 6     self.executor = concurrent.futures.ThreadPoolExecutor(max_workers=len(self.cluster))
```

3.1 Node Types

As previously stated, there are three node types: leader, follower and candidate (code ??). In this section we are going to show their characteristics and similarities. Note that all nodes have a timer: it is randomized for each of them and has been implemented by extending *threading.Timer*, thus making it thread-safe (code ??)

Listing 7. Node modes

```
3   1 class Raft(SimpleXMLRPCServer):  
4     2     class Mode(Enum):  
5       3         LEADER = 1  
6       4         CANDIDATE = 2  
7       5         FOLLOWER = 3
```

```

469   6
470   7     def __init__(self,
471   8         # ...
472   9         mode: Mode = Mode.FOLLOWER,
473  10        )

```

474

475

476

477

478

Listing 8. Threadsafe looping timer

```

479  1 class LoopTimer(Timer):
480  2     def __init__(self, interval, function, args=None, kwawrgs=None):
481  3         Timer.__init__(self, interval, function, args, kwawrgs)
482  4         self.was_reset : bool = False
483  5         # ...
484  6
485  7 class Raft(SimpleXMLRPCServer):
486  8     def __init__(self,
487  9         # ...
488 10        )
489 11     # start timer
490 12     self.timer = LoopTimer(timeout, self.on_timeout)
491 13     self.timer.start()

```

492

493

3.1.1 *Leader Node.* The algorithm revolves around, and requires the existence of, one and only one leader node, whose job is to synchronize all servers' logs to ensure data consistency. It does so by replicating its own log on all followers (the non-leader nodes) by sending new or, if needed, old entries via remote procedure calls.

To make sure all nodes believe the leader's alive, periodically sends an empty remote procedure call called *heartbeat* (every 150-300ms).

500

3.1.2 *Follower Node.* All nodes, except for the leader, perform as followers. They are not allowed to replicate their own log, and they have to forward any request to the leader.

To make sure the cluster never remains without a leader, every follower has an election timeout (between 150ms and 300ms) which resets every time an RPC from the leader is received. If it times out, the follower changes its state to *candidate*, increments its current term and starts a leader election.

Followers become candidates in another scenario: whenever they receive an entry from the leader, they compare it with their own last log entry. If the leader's term is smaller, it is out of date and a new election is started.

510

3.1.3 *Candidate Node.* When a follower's election timeout times out, it becomes a candidate, increments its own term and starts an election. Votes for itself and then waits for one of two outcomes: wins, thus becoming a new leader, or loses (either another leader gets elected or the old one manifests itself) thus reverting back to being a follower.

514

515

3.2 Log

517

As stated, the leader's job is to accept requests (in our specific case they are player inputs) and then forward them to the followers. Let's talk about the structure of the log.

519

Manuscript submitted to ACM

521 The log is basically a *list* (or an *array*) of entries, where *entry* is an element that encapsulates data (like an integer or
 522 a string), has an index (unique for each entry) and the term of its creation (figure ??). We defined entries as *Data Classes*
 523 ¹⁴ (decorators that simulate C's structures) as seen in code ??.

525

526

Listing 9. Dataclass Entry definition

```
527 1 @dataclass
528 2 class Entry:
529 3     term: int
530 4     index: int
531 5     command: str
```

532

533

534

535

536

537

538

539

540

541

542

543

544

545

546

547

548

549

images/logStructure.png

550

551

552

553

554

555

556

557

558

559

560

561

562

563

564

565

566

567

Fig. 5. Raft's log is fundamentally an array made of entries

568

569

570

¹⁴Data Classes module provides a decorator and functions for automatically adding generated special methods to user-defined classes: <https://docs.python.org/3/library/dataclasses.html>

571

572

3.3 Log Replication and Overwriting

All log propagation revolves around one remote procedure call named *append_entries_rpc*, which the leader calls on a list of server proxies that connect it to the followers. On their end, each follower calls the RPC on a proxy of the leader. It must be registered in the server to be callable, as seen in listing ??.

Listing 10. Register, thus making it callable, the remote procedure call *append_entries_rpc*

```

573
574
575
576
577
578
579
580     def handle_server():
581         with Raft(...) as server:
582             def append_entries_rpc(entries, term, commit_index, all_log):
583                 ...
584                 server.register_function(append_entries_rpc)      # makes function callable on the other side
585                 server.serve_forever()                         # keeps server alive
586

```

3.3.1 *Leader Propagates Entries*. The leader (each node as a matter of fact) periodically checks whether there are new commands to propagate (always stored in queue *pygame_commands*, more details in section ??), by overriding *SimpleXMLRPCServer*'s method *service_actions* (listing ??).

Then, it translates them into entries by giving each of them the current term and a log index that starts from *lastLogEntry(index)* + 1 and increases by one for each entry. To clarify: if *lastLogEntry(index)* = 7 and we have three new commands, their indices will respectively be eight (8), nine (9) and ten (10). The translation can be seen at listing ??.

At this point, it propagates *new_entries* to the whole cluster, updating the commit index (necessary for applying log to state) as soon as propagations are successful on at least half of the cluster, like so: *commitIndex* = *lastNewEntry(index)*.

What happens if the *append entries* gets rejected? The leader adds to *new_entries* its own last log entry: *new_entries* = *lastLogEntry* + *new_entries* (figure ??). Then, it repeats the propagation procedure, for each reject a new *last log entry* gets added, progressively traversing the log backwards. If, at a certain point, *new_entries* == *allLog* + *new_entries* (i.e., all leader's log gets propagated) the flag *all_log* is set to *True*.

Since every server may reject or accept different sets of entries, depending on their own local log, every propagation must be "local" for each follower.

The flow of execution for the log propagation is: *Raft: service_actions* → *Raft: propagate_entries* → *propagate_entries :encapsulates_proxy: append_entries_rpc*. The last one gets called as many times as needed on every single follower.

Of course, all propagation happens concurrently using a *ThreadPoolExecutor*, and the code for entries propagation (leader's side) can be seen at listing ??.

Listing 11. Periodically checks whether there are new commands

```

610
611
612     def service_actions(self):
613         if time.time() - self.countdown >= .005:    # RUN method of the server, override
614             global pygame_commands
615
616             if not pygame_commands.empty():           # do actions every .005 seconds
617                 self.propagate_entries()

```

Listing 12. Translates commands into new entries

```

618
619
620     def propagate_entries():
621     # ...
622     while not pygame_commands.empty():
623         command = pygame_commands.get()

```

```

625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640     images/newEntries.png
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659

```

Fig. 6. New entries with the last log's entry pushed to the top of the list

```

660     5         log_index += 1                      # lastLogEntry(index) + 1
661     6         self.new_entries.append(Raft.Entry(   # append to support list 'new_entries'
662     7             term= self.term,
663     8             index= log_index,
664     9             command= command
665    10        ))
666
667
668

```

Listing 13. Leader propagation procedure, for complete code refer to project's repository

```

669 1 def propagate_entries(self):
670 2     # ...
671 3     if self.log:
672 4         entries: list[Raft.Entry] = []
673 5         entries.append(self.log[-1])
674 6         entries.extend(self.new_entries)
675 7         log_iterator: int = -2           # log_iterator soft resets for each follower
676 8     else:

```

```

677   9     entries: list[Raft.Entry] = self.new_entries
678  10    log_iterator: int = -1
679  11
680  12    # ...
681  13    # inner function necessary for concurrent execution
682  14    def encapsulate_proxy(self, follower, entries, log_iterator):
683  15        # ...
684  16        with xmlrpclib.ServerProxy(complete_url, allow_none=True) as proxy:
685  17            while not propagation_successful:
686  18                # send new entries (local for each follower)
687  19                # ...
688  20                result = proxy.append_entries_rpc(entries, self.term, self.commit_index, all_log)
689  21                if result[0] == False:
690  22                    # add another entry from self.log to new entries
691  23                    entries = [self.log[log_iterator]] + entries
692  24                    log_iterator -= 1
693  25                elif result[0] == True:
694  26                    propagation_successful = True
695  27
696  28    return propagation_successful # to propagate_entries, make propagation counter increase
697  29
698  30    results = []
699  31
700  32    # fires RPCs concurrently using ThreadPoolExecutor
701  33    future_result = {           # clever python syntax trick
702  34        self.executor.submit(
703  35            encapsulate_proxy,      # function
704  36            self,                  # function's parameter
705  37            follower,              # function's parameter
706  38            entries,               # function's parameter
707  39            log_iterator           # function's parameter
708  40        ): follower for follower in self.cluster}
709  41    for future in concurrent.futures.as_completed(future_result):
710  42        # results of RPCs
711  43        data = future.result()
712  44        results.append(data)
713  45
714  46    # finally counts if propagation was successful enough
715  47    if results.count(True) >= len(self.cluster) / 2:
716  48        self.log.extend(self.new_entries)          # add new entries to log
717  49        self.new_entries.clear()                  # clear new entries list
718  50        self.commit_index = self.log[-1].index    # ensure log gets eventually applied
719  51
720  52    else:
721  53        # new entries are not cleared, so they will be propagated again
722
723
724
725
726
727
728

```

3.3.2 *Follower Receives Entries*. When a follower receives an *append entries* request from the leader, first checks whether leader is up to date. If it's not, i.e., $leaderTerm < followerTerm$, rejects by answering with the tuple $(False, followerTerm)$. In this context, *answering* is done via the remote procedure call's return value.

On the other hand, if the leader's term is equal or greater than its own (i.e., $leaderTerm \geq followerTerm$), the follower updates its commit index and, if $leaderEntries \neq \emptyset$, checks the *all_log* flag. If it's *True*, it clears all its own log to overwrite it with the leader's (fundamental to log forcing, listing ??). Otherwise ($all_log \neq True$), the leader did not

729 send all its log, so the follower searches through its own log for an entry equal to the leader's previous one (i.e., the
 730 entry preceding the new ones). Let's make an example:
 731

- 732 • Leader's log = [1, 2, 3, 4, 5];
- 733 • Leader's new entries = [6, 7];
- 734 • Thus leader's prev = [5].

735 If it finds an entry equal to leader's previous (i.e., $followerLog(someEntry) == leaderPrev$), deletes all log entries
 736 that follow it and appends the new ones, otherwise ($\#(followerLog(someEntry) == leaderPrev)$) rejects the request.
 737 Since the leader, when faced with a reject, adds a new *prev* and keeps repeating the send until it comprises all its log, at
 738 a certain point the follower will be forced to overwrite all its log, thus making it equal to the leader's. This overwriting
 739 is called *log forcing* and ensures that all logs are equal to the leader's.
 740

741 The code can be seen at listing ?? (for the complete one refer to the repository).

742
 743 Listing 14. Follower clears its own log to overwrite it with the leader's

```
744 1 if all_log == True:
  2     server.log.clear() # if leader sent all its log, clear and rewrite log (leader's log forcing)
  745
```

746 Listing 15. Follower search in its own log for an entry equal to leader's prev

```
747 1 if commit_index is not None:
  2     server.commit_index = commit_index # update commit index
  748
  3
  4 if entries is not None:      # not an heartbeat
  5     if all_log == True:        # complete overwrite
  6         server.log.clear()
  7
  8     if server.log: # if follower's log not empty search for an entry equal to leader's prev
  9         entry_log_index: int | None = None           # save its log index (!= entry index)
 10
 11         for i, my_entry in enumerate(server.log):
 12             if (my_entry.index == entries[0].index
 13                 and my_entry.term == entries[0].term):
 14                 entry_log_index = i
 15                 break # no need to search further
 16
 17             # here entry_log_index == (position of entry equal to leader.prev) | None
 18
 19
 20         if entry_log_index is None:          # entry equal to leader's prev not found
 21             return(False, server.term)      # rejects
 22
 23
 24         del server.log[(entry_log_index ):] # delete all log following leader prev
 25
 26
 27         server.log.extend(entries) # append new entries
 749
```

750 3.3.3 *Follower Sends Entries*. Since every server is a Raftian node with a game instance and thus player inputs, followers
 751 have their own Pygame commands to propagate. Just like the leader, in their *service_actions* function they periodically
 752 check whether there are new commands to propagate and call *propagate_entries* accordingly. Then, they translate
 753 all Pygame commands into entries (same code as listing ??) and propagate them to the leader via *append_entries_rpc*.
 754 Nothing else.

755 As previously stated, followers are *passive*, meaning they do not apply their own player inputs when they register
 756 them, but only after the leader propagates them back to the whole cluster.

781 3.3.4 *Leader Receives Entries.* The leader does very little when receives entries from the followers: it just puts them
782 into its own *pygame_commands* queue. They will get processed and propagated eventually, as stated in section ??.
783

3.4 Apply Log to State

Let's first explain two key attributes: *commit index* and *last applied*. Both of these represent an index, but the former is the highest-index entry successfully propagated in the cluster, while the latter is the highest-index entry already applied to state.

Every node, whether leader or follower, applies entries to state in the same way: inside their function *service_actions* they periodically check if there is a discrepancy between *commit index* and *last applied* attributes (i.e., $\text{commit_index} > \text{last_applied}$). Then, starting from the last applied entry, they apply to state all successive entries up to and including the one with the same index as *commit_index*, updating *last_applied* as they go. To clarify: servers apply all entries between $\log(\text{entry.index} == \text{last_applied})$ and $\log(\text{entry.index} == \text{commit_index})$ as shown in figure ??.

To apply entries in our context means that they get appended to the queue `raft_orders`. The code can be seen at listing ?? (for the complete source refer to the repository)

Listing 16. All nodes apply entries to state based on *commit_index*

```
800
801     def service_actions(self):
802         # ...
803
804         if self.commit_index is not None and self.commit_index > self.last_applied:
805             global raft_orders # applying means appending entries to this queue
806
807             #...
808
809             last_applied_log_position: int = -1
810             for i, my_entry in enumerate(self.log):
811                 if (my_entry.index == self.last_applied):
812                     last_applied_log_position = i
813                     break # found log position of last applied entry
814
815             log_iterator = last_applied_log_position + 1      # improves code clarity
816
817             while self.last_applied != self.commit_index:
818                 raft_orders.put(self.log[log_iterator])
819                 self.last_applied = self.log[log_iterator].index
820                 log_iterator = log_iterator + 1
821
822             # here self.last_applied == self.commit_index
```

3.5 Log Compaction

This functionality, as well as the following ones, have not been implemented due to time constraints. We decided to include them anyway since they were still the product of careful consideration and could prove useful in future implementations.

Log compaction, also called *snapshot*, is a way to clear servers' logs and save them in persistent memory, necessary in long-running or message-heavy applications. Every node autonomously decides when to do it.

The idea is as follows: inside their function *service_actions*, servers check whether their log is larger than a certain size (to be determined) and then call a *snapshot* method accordingly, which saves all entries in a JSON file progressively deleting them from the log, from the first one up to (but excluding) the *last applied*.

```
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848     images/applyToState.png
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
```

Fig. 7. All entries between *last_applied* and *commit_index* (included) get applied to state

An alternative interpretation, closer to Ongaro and Ousterhout's idea, is saving the state of the cluster. When applied to our context, we can imagine a JSON file that describes all servers' state, by saving for each of them *id*, *url*, *port*, and *hp* values. *Index* and *term* of the last snapshotted entry should also be saved, as well as the current configuration (which will be mentioned in section ??). An example of the JSON structure can be seen at listing ??.

This method requires more pre-processing to generate the JSON, and more post-processing to later check log's correctness, but is also more compact, which helps when the leader calls *install snapshot*, a remote procedure call that can sometimes be needed to bring up to speed new or outdated servers by sending them the leader's snapshot.

It goes without saying, but whichever method is ultimately chosen, every new snapshot must comprise all information contained in the old ones.

Listing 17. The JSON for a snapshot that saves cluster's state would look something like this

```
880
881   1 {
882   2     servers: [
883   3       {
```

```

885   4         id : 1,
886   5         url : "localhost",
887   6         port : 8000,
888   7         hp : 70
889   8     },
890   9     {}, {}
900 10   ],
911 11   lastIndex : 11,
922 12   lastTerm : 3,
933 13   lastConfig : 8
944 14 }
955
966

```

3.6 Leader Election

As aforementioned, the leader is fundamental to grant consistency. To make the protocol fault tolerant, it can dynamically change over time via a distributed election: whenever a follower finds out that the leader is either outdated or missing (i.e., internal follower's timer times out before receiving any call from it), said follower starts an election. It changes its internal state to *candidate*, increases its own term by one, votes for itself, and then propagates to the whole cluster a specialized remote procedure call named *request_vote_rpc* (code, removed in the final version, at listing ??). Votes are given on a first-come-first-served basis, and to prevent split votes each server's election timeout is randomized between 150ms and 300ms at the start of every election. This ensures that in most cases only one server will be candidate at a time.

At this point there are two possible outcomes: more than half of the cluster votes for the candidate (which we will call "*A*"), that therefore becomes leader and propagates a heartbeat to the whole cluster, or another candidate (which we will call "*B*") is more up-to-date (i.e., *B*'s *term* is greater than *A*'s or equal but with a greater *lastIndex*). In this last case, candidate *A* reverts back to follower and votes for *B*. The pseudocode for all the above can be seen at listing ??.

One last eventuality is that the old leader manifests itself. In this case, if the old one is equally or more up-to-date than the new one (both term and last index count), the latter reverts back to follower and the preceding monarch gets reinstated.

Listing 18. Pseudocode for *request_vote_rpc*

```

919 1 def request_vote_rpc(...):
920 2     # if candidate less up to date -> reject
921 3     if self.term > candidate_term:
922 4         return (self.term, False)
923 5
924 6     # if a candidate already exists
925 7     if self.voted_for is not None and not candidate_id:
926 8         return (self.term, False)
927 9
92810     # vote for candidate
92911     self.voted_for = candidate_id
93012     return (self.term, True)
93113 #...
93214 server.register_function(request_vote_rpc)

```

Listing 19. Pseudocode for *to_candidate*, gets fired on *election timer* timeout

```

933
934 1 def to_candidate(self):
935
936 Manuscript submitted to ACM

```

```

937 2     self.mode = Raft.Mode.CANDIDATE
938 3     self.term += 1
939 4     self.voted_for = self.id
940 5
941 6     self.timer.reset() # reset election timer
942 7
943 8     for server in self.cluster:
944 9         server.request_vote_rpc(...)
945 10        count votes
946 11
947 12     if some_return.more_up_to_date:
948 13         self.mode = Raft.Mode.FOLLOWER
949 14         self.voted_for = some_return.id
950 15
951 16     if votes > len(self.cluster) / 2:
952 17         self.to_leader()    # handles mode change and heartbeat

```

3.7 What is Missing

There is one last functionality discussed by Ongaro and Ousterhout, which is gracefully managing changes in the cluster's members by leveraging a configuration attribute and keeping multiple configurations alive simultaneously for a certain period of time (figure ??). We never intended to include it due to time constraints, therefore there is nothing we can add beyond the original work.

Another concern, which was not considered in the Raft paper, is faults caused by bad actors that purposely send malicious information. This is a real problem for our use case, since players want to win and are therefore incentivized to act maliciously by cheating (a practice so widespread that has created its own multimillion-dollar market [?]).

The implementation of Byzantine fault tolerance was beyond the scope of this work. Therefore, we refer the reader to some example works for further details: "*A Raft Algorithm with Byzantine Fault-Tolerant Performance*" by Xir and Liu [?], and "*VSSB-Raft: A Secure and Efficient Zero Trust Consensus Algorithm for Blockchain*" by Tian et al. [?].

4 RAFTIAN NODE ARCHITECTURE

In the previous section (??), we explained in detail how nodes communicate with each other and handle their log. Now we will explain what actually happens inside a node, i.e., the architecture of a single node of the application, comprising of both a server and a game instance. Before showcasing it, we will briefly explain how Pygame works, thus taking the opportunity to present the user interface.

4.1 Pygame

Pygame's approach is very straightforward: first comes the declaration and set up of all the graphical components, such as game window, fonts, colors, variables and constants. Every element gets positioned on the main window by coordinates (x,y) , where $(0,0)$ is the top left corner. Most items are made of two fundamental Pygame classes: *Rect*, which creates non-graphical objects that expose many useful methods, for example to position, move, and resize themselves, or to detect collisions and mouse clicks, and *Surface*, which is the most basic graphical component that has dimensions and can be drawn upon. Often we want to bind, thus constrain, surfaces with rects so that we use the latter for spatial operations. One last fundamental is the *blit* function, a method that draws one image onto another or, to be precise, that draws a source Surface onto the object Surface that calls it. We can give it an optional argument to specify a

```

989
990
991
992
993
994
995
996
997
998
999
1000 images/raftMembershipChange.png
1001
1002
1003
1004
1005
1006
1007
1008
1009
1010
1011
1012
1013 Fig. 8. Cluster goes through a hybrid configuration to pass from the old to the new one. Source: Raft paper [? ]
1014
1015
1016 drawing destination, either with coordinates or a rect. To clarify: baseSurface.blit(sourceSurface, destination) draws
1017 sourceSurface onto baseSurface at the coordinates specified by destination. An example of all the above can be seen at
1018 listing ??.
```

Pygame, being low-level in nature, is very flexible and allows us to do pretty much whatever we want. For example, we defined our players as dataclasses that encapsulate both the players' data (like *id* or *health points*) and their Rect and Surface objects, as in listing ??.

Finally, Pygame gives us direct access to the game loop, which is implemented as nothing more than a *while loop*. In it, we can process player inputs and refresh the screen, dynamically changing what is displayed. In short, we manage everything that happens while the game is running. In listing ?? we can see two types of player input, one for quitting the game and a left-mouse click, the latter of which causes a refresh of the header. Specifically, if *Player 2* gets clicked, the header will display "*Player 2 pressed*", reverting back to its original state after a couple of seconds. The last command, *clock.tick(fps)*, allows us to limit the framerate, effectively slowing down or speeding up the game engine itself by constraining the amount of times per second the game loop repeats itself.

Listing 20. Pygame base components

```

1034
1035 1 pygame.init() # starts pygame
1036 2 GREY = (125, 125, 125) # define a color
1037 3 DISPLAY = pygame.display.set_mode((1000, 1200)) # creates game window 1000x1200 pixels in resolution
1038 4 clock = pygame.time.Clock() # necessary to mange fps
1039 5 font = pygame.font.Font(None, 60) # creates default font
1040 6 toptext = font.render("Top Text", False, BLACK) # header text
```

```

1041 7 rect_header = pygame.Rect(0, 0, 1000, 100)      # creates rect for header
1042 8 header = pygame.Surface((1000, 100))        # creates surface for header
1043 9 header.fill(WHITE)                          # draw on surface
1044 10
1045 11 DISPLAY.blit(header, rect_header)          # draw on DISPLAY the header surface
1046 12                                         # position is given by rect_header
1047 13 #...
1048 14 # draw text on coordinates
1049 15 DISPLAY.blit(toptext, (rect_header.centerx - xoffset, rect_header.centery - yoffset))
1050
1051
1052

```

Listing 21. Players defined as dataclasses that encapsulate Pygame elements

```

1053 1 @dataclass
1054 2 class Player:
1055 3     id: int
1056 4     hp: int
1057 5     rc: pygame.Rect    # represents player position and size
1058 6     ui: pygame.Surface # exposes UI of the player e.g., colour
1059 7
1060 8     player1 = Player(
1061 9         id=1,
1062 10         hp=100,
1063 11         rc=pygame.Rect(585, 685, 80, 80), # x0, y0, width, height
1064 12         ui=pygame.Surface((80,80))
1065 13     )
1066 14     player1.ui.fill(RED)                  # colour player red
1067 15     DISPLAY.blit(player1.ui, player1.rc) # draw on display via rect
1068

```

Listing 22. All interactions and frame-by-frame rendering happen in the game loop

```

1069
1070 1 while True: # game loop
1071 2     for event in pygame.event.get():    # process player inputs
1072 3         if event.type == pygame.QUIT:
1073 4             pygame.quit()
1074 5
1075 6         if event.type == pygame.MOUSEBUTTONDOWN and event.button == 1: # left mouse button click
1076 7
1077 8             pos = pygame.mouse.get_pos() # gets mouse position
1078 9
1079 10             if player1.rc.collidepoint(pos): # rect allows us to detect collisions
1080 11                 toptext = font.render(f"Player {player1.id} pressed", False, BLACK)
1081 12
1082 13                 DISPLAY.blit(header, rect_header) # erase previous text
1083 14                 DISPLAY.blit(toptext, (rect_header.centerx - xoffset, rect_header.centery - yoffset))
1084 15
1085 16             pygame.display.flip() # refresh on-screen display
1086 17             clock.tick(60)       # limits framerate
1087

```

4.2 Raftian User Interface

Figure ?? shows different phases of a normal Raftian's game session. Specifically, they demonstrate how the interface changes when the player repeatedly clicks on (thus *attack*) Player 3, the blue one on the top left (players are represented

1093 as four coloured squares on the board). Players' colours become progressively desaturated as their health points decrease,
 1094 by modifying their alpha channels as shown in listing ???. When a player is dead, its colour changes to a darker shade.
 1095

1096 In the header an attack message gets written, reverting back after half a second. Said message changes depending
 1097 whether the player is still alive; if not, damage is ignored.

1098

1099

Listing 23. Whenever a player gets damaged, its colour gets desaturated

```
1100 1 for player in players:
1101 2     if player.id == order.command and player.hp > 0:
1102 3         player.hp -= 30 # apply damage to player:
1103 4
1104 5     if player.hp < 90 and player.hp >= 60:
1105 6         player.ui.set_alpha(190)
1106 7         DISPLAY.blit(player_UI_cleaner, player.rc) # clean player UI
1107 8         DISPLAY.blit(player.ui, player.rc)          # redraw player UI
```

1108

1109

1110

1111

1112

1113

1114

1115

1116

1117

1118

1119

1120

1121

1122

images/raftian1.png

images/raftian2.png

images/raftian3.png

(a) Game starts

(b) Player 3 is attacked

(c) Header reverts to default

images/raftian4.png

images/raftian5.png

images/raftian7.png

(d) Player 3 is left with forty HP

(e) Player 3 is left with ten HP

(f) Player 3 cannot be attacked again

1135

1136

Fig. 9. Different phases of a normal Raftian's game session. Player 3 keeps getting damaged until it dies

1138

1139

4.3 Raftian Node Architecture

1140

1141

The architecture of a Raftian node can be seen at figure ???. Let's explain it: first of all, the game loop and the server are encapsulated in two functions to be handed over to two different threads, enabling concurrent execution (listing ???).

1142 Manuscript submitted to ACM

1143

1144

1145 Whenever a player clicks on (i.e., attacks) one of the four players, the game engine does not apply damage immediately.
 1146 Instead, it generates a *command* which represent, if we want, the *intention* of attacking said player. This *command* is
 1147 thus appended to a queue called *pygame_commands*, one of the two synchronized FIFO queues¹⁵ necessary to allow
 1148 communication between server and Pygame's threads (listing ??). Both are instances of Python's standard library queue
 1149 module¹⁶, which implements thread-safe, multi-producer, multi-consumer queues.

1151 At this point, Pygame does not concern itself anymore with said user input. The server, by itself, periodically checks
 1152 the *pygame_commands* queue (as in listing ??) and, when not empty, removes elements from it (as in listing ??) and
 1153 propagates them as entries to the leader (or to the whole cluster if said server *is* the leader, as in listing ??).

1155 Then, the leader propagates the received commands to the whole cluster, which we will now call *orders*. Each server
 1156 adds received orders to its own log, as explained in section ??, so that they can later be appended to the *raft_orders*
 1157 queue when entries get applied to state (as in section ??). This way, the original user input gets propagated back to the
 1158 server that generated it in the first place.

1159 Finally, Pygame checks (periodically) the *raft_orders* queue for orders. When it finds them, it removes them from the
 1160 queue and updates the user interface accordingly (an example can be seen at listing ??).

1162 The whole idea is to keep server and game engine as separated as possible: the former reads commands, propagates
 1163 them and writes received orders, the latter reads orders, updates the UI, and writes commands, following a unidirectional
 1164 cyclic communication pattern.

1167 Listing 24. Start both Pygame and server's threads

```
1169 1 def handle_pygame():
1170 2     pygame.init()
1171 3     #...
1172 4     While True:
1173 5         #...
1174 6 def handle_server():
1175 7     with Raft(...) as server:
1176 8         #...
1177 9     server.serve_forever()
117810
117911 server_thread = threading.Thread(target=handle_server)
118012 server_thread.start()
118113 pygame_thread = threading.Thread(target=handle_pygame)
118214 pygame_thread.start()
```

1183 Listing 25. Queues for commands and orders, they allow inter-thread communication

```
1185 1 # user inputs through Pygame which writes them here
1186 2 # Raft reads them and propagates them to the cluster
1187 3 pygame_commands = Queue()
1188 4
1189 5 # commands that have been applied to state are written here by Raft
1190 6 # Pygame reads them and updates UI accordingly
1191 7 raft_orders = Queue()
```

1192
 1193 ¹⁵FIFO, or First-In-First-Out, is a method for organizing the manipulation of a data structure, often data buffers, where the oldest data inserted is the first
 1194 that gets processed, making it work in a sense like a pipeline

1195 ¹⁶Python's queue, a synchronized queue class: <https://docs.python.org/3/library/queue.html>

1197 Listing 26. Pygame periodically checks whether there are new orders and updates the UI accordingly

```

1198 1 while True: # Pygame's main loop
1199 2     ...
1200 3     while not raft_orders.empty():
1201 4         order: Raft.Entry = raft_orders.get()
1202 5
1203 6         for player in players:
1204 7             if player.id == order.command and player.hp > 0:
1205 8                 player.hp -= 30 # apply damage to player:
1206 9             ...

```

1207

1208

1209

1210 4.4 Evaluation and Results

1211 While we did not have enough time to implement evaluation procedures for measuring things like response time or
 1212 network latency, the game has been thoroughly tested. Responsiveness is good, with no noticeable input latency and
 1213 a solid framerate way above sixty frames per second (which is still the preferred limit), and both log replication and
 1214 overwriting functionalities have been confirmed working as intended. The proof of this fact can be observed in the
 1215 logs at <https://github.com/mhetacc/RuntimesConcurrencyDistribution/tree/main/logs>, specifically by comparing logs
 1216 among folders *bob1*, *bob2*, *bob3*, *bob4*, and *raftian*: all logs (meaning *Raft logs*) contain the same entries, in the same
 1217 order, between all players.

1218

1219

1220 5 REFLECTION

1221 Let's now discuss problems, potential future expansions and learning outcomes of this project.

1222

1223

1224 5.1 Self-Assessment

1225 Using *XML-RPC* and *threading* libraries proved to be sub-optimal: the former has a very contrived syntax and makes
 1226 writing procedure calls a bit unintuitive, since forces the programmer to think in the "opposite direction". When writing
 1227 a remote procedure call (i.e., those functions that get registered by *register_function()*) is important to keep in mind that
 1228 they are going to be used by the caller and not by sender (in whom code block they are written in).

1229

1230

1231

1232

1233

1234

1235

Code could be less coupled: both server and game loop reside in the same file, and a lot of components are either internal classes or nested functions. Moreover, both command and order queues are global variables, which is generally a practice to be avoided.

1236

1237

1238

On the other hand, code is well documented and as understandable as possible, even though following the flow of, for example, an input propagation requires jumping through it many times.

1239

1240

5.2 Future Works

1241

1242

1243

1244

1245

1246

1247

Apart from the two features already discussed (leader election and log compaction), future expansions could implement cluster's membership change and a Byzantine fault-tolerant version of Raft. Adding new game functionalities, thus command types, should be easy since they can be propagated by the existing infrastructure, and the same is true for adding new players: provided that a new, bigger, user interface gets created, changing cluster's size should, in our testing, work without any issues.

1248

```
1249  
1250  
1251  
1252  
1253  
1254  
1255  
1256  
1257  
1258  
1259  
1260  
1261  
1262  
1263  
1264  
1265  
1266  
1267  
1268 images/nodeArchitecture.png  
1269  
1270  
1271  
1272  
1273  
1274  
1275  
1276  
1277  
1278  
1279  
1280  
1281  
1282  
1283  
1284  
1285  
1286  
1287  
1288
```

Fig. 10. Raftian node architecture

5.3 Learning Outcomes

We started this project by having very limited Python competencies, having never written concurrent programming, never touched a game engine and never worked with network transmission protocols. All in all, we learned all of the above, in some cases going so far as trying different alternative solutions (we implemented Raft nodes mockups with threading, multiprocessing and asyncio libraries), making this project an invaluable learning experience.